

Odessa: Enabling Interactive Perception Applications on Mobile Devices

Moo-Ryong Ra*, Anmol Sheth⁺,
Lily Mummert^x, Padmanabhan Pillai',
David Wetherall^o, Ramesh Govindan*

**USC ENL, ⁺Technicolor, ^xGoogle, 'Intel,
^oUniversity of Washington*

Emerging Mobile Perception Applications



Dual-Core CPU

Computation



Cloud Infrastructure

Communication

Activity
Recognition

Health, Traffic
Monitoring

Location-Based
Service

Participatory
Sensing

Sensing Applications





Vision-based Interactive Mobile Perception Applications

Face
Recognition

Object and Pose
Recognition

Gesture
Recognition



Motivation



Problem



Measurement



Design



Evaluation

Common Characteristics

Interactive

- Crisp response time (10 ms ~ 200 ms)

High Data-Rate

- Processing video data of 30 fps

Compute Intensive

- Computer Vision based algorithms



Enabling Mobile Interactive Perception


Performance

Throughput 

Makespan 

Application	Throughput	Makespan
Face Recognition	2.50 fps	2.09 s
Object and Pose Recognition	0.09 fps	15.8 s
Gesture Recognition	0.42 fps	2.54 s

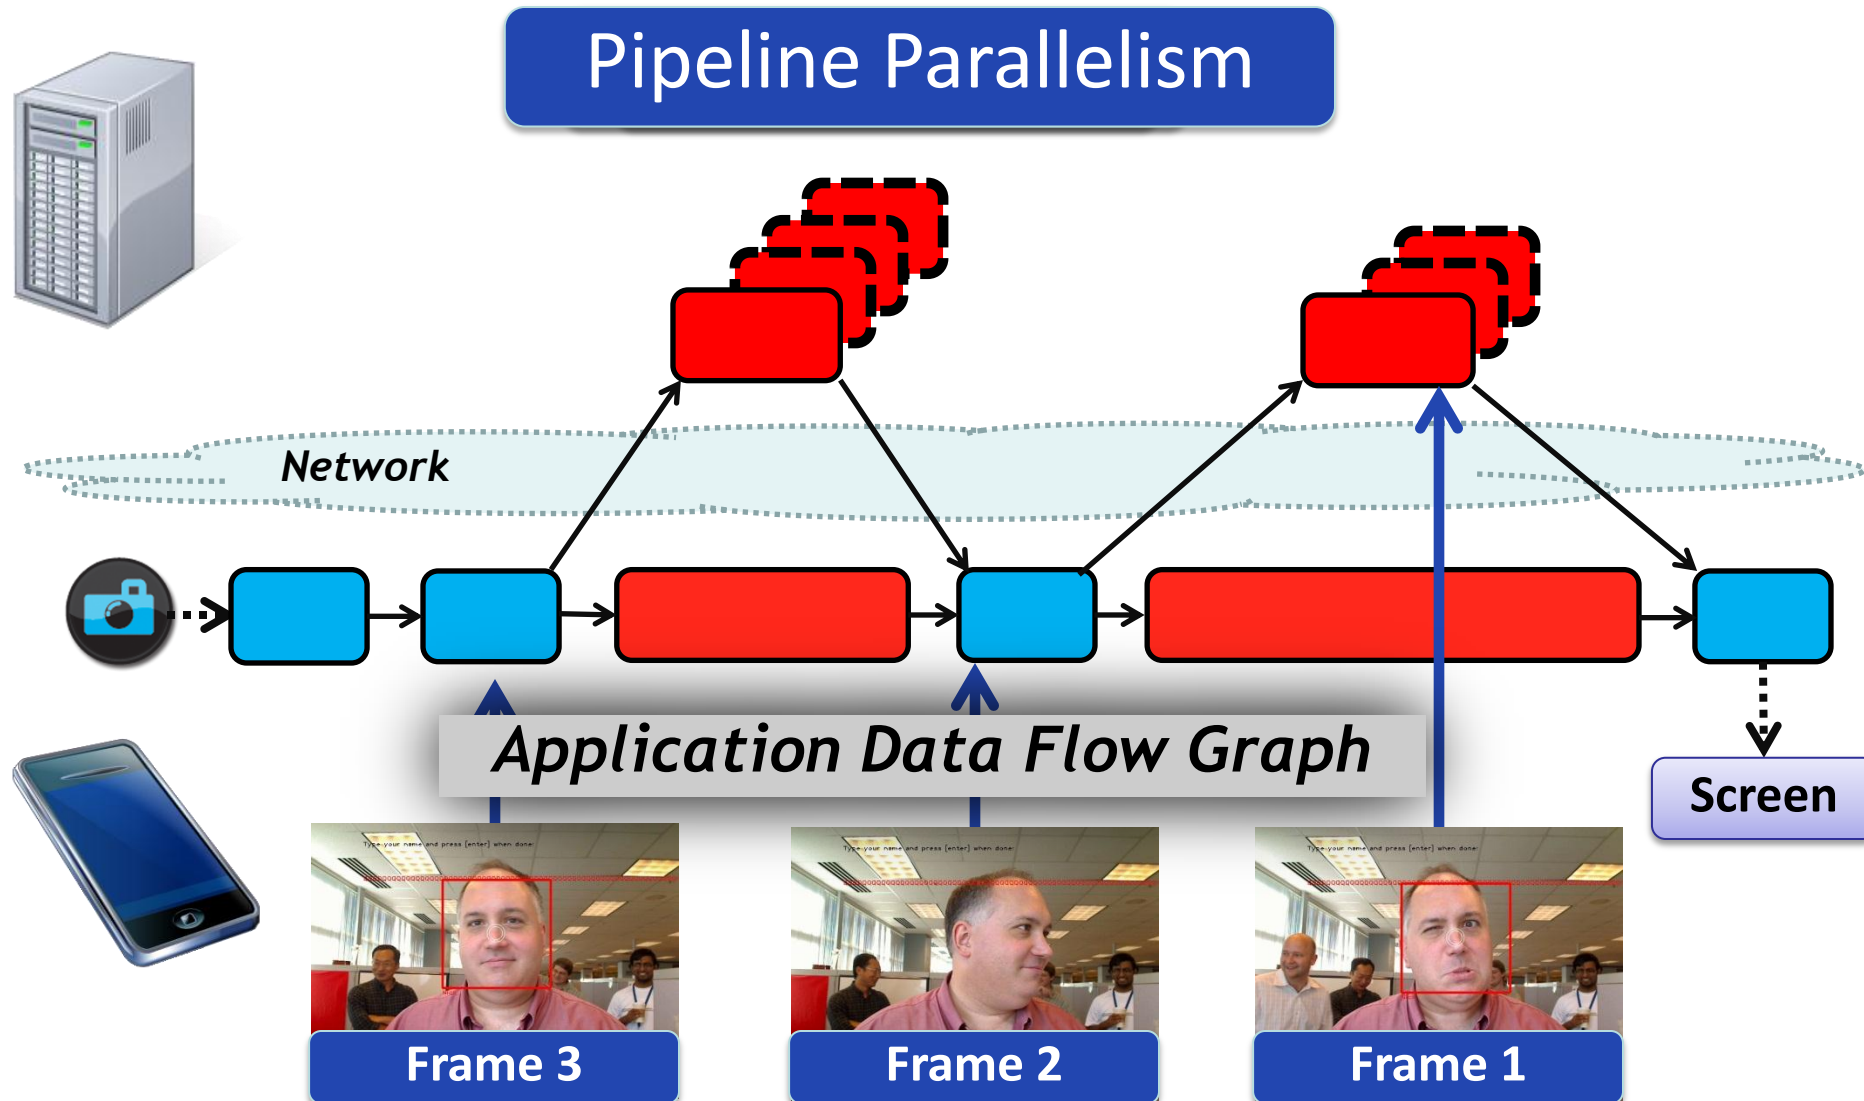
All running locally on mobile device



Video of 1 fps



Two Speed-up Techniques



Main Focus

Data Flow Structure



Offloading



Parallelism

System Support



Enable Mobile Interactive Perception Application



Contributions

What factors impact offloading and parallelism?

Measurement

How do we improve
throughput and makespan simultaneously?

Odessa Design

How much benefits can we get?

Evaluation



Measurement

Input Data Variability

Varying Capabilities of Mobile Platform


Network Performance

Effects of Parallelism




Lesson I : Input Variability

Object and Pose Recognition



The system should adapt to the variability at runtime




Impact of **input variability**

Lesson II: Effects of Data Parallelism

Object and Pose Recognition

# of Threads	Thread 1	Thread 2	Thread 3
<p>The level of data parallelism affects accuracy and performance.</p>			



**Input
Complexity**

**Segmentation
Method**



Summary: Major Lessons

Offloading decisions must be made in an adaptive way.

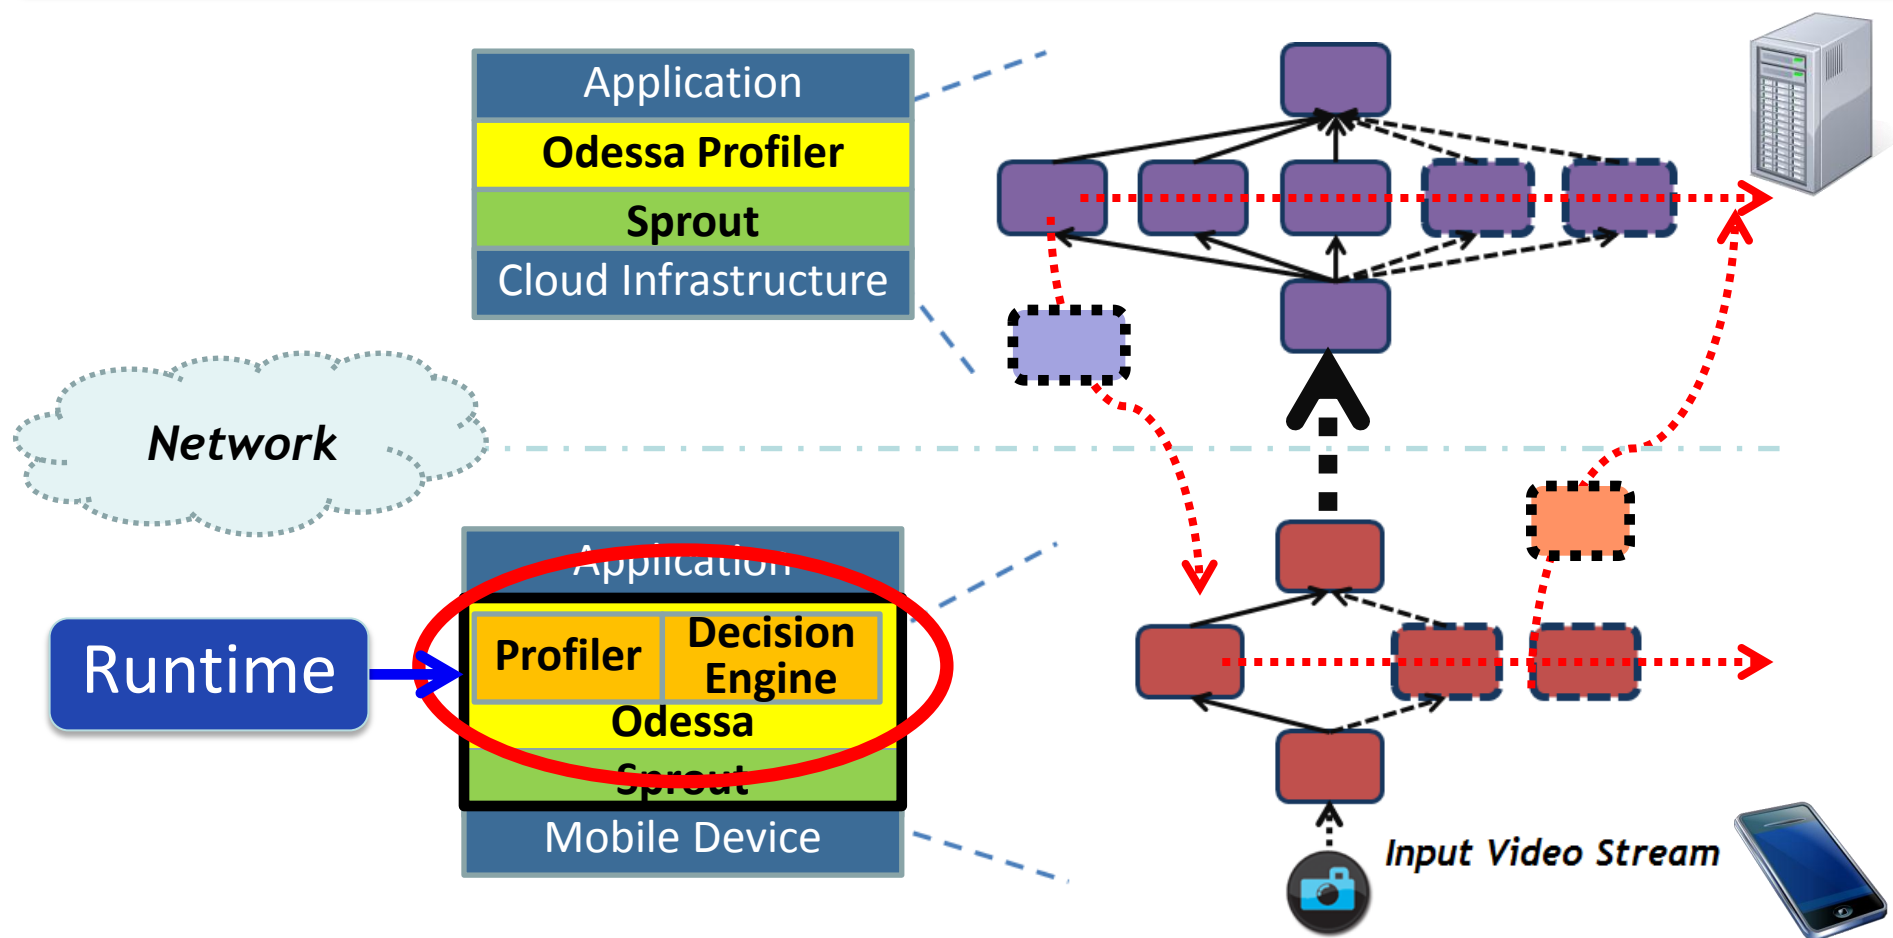
The level of data parallelism cannot be determined a priori.

A static choice of pipeline parallelism can cause sub-optimal performance.




Odessa

Offloading DDecision System for Streaming Applications



Incremental Decision Making Process

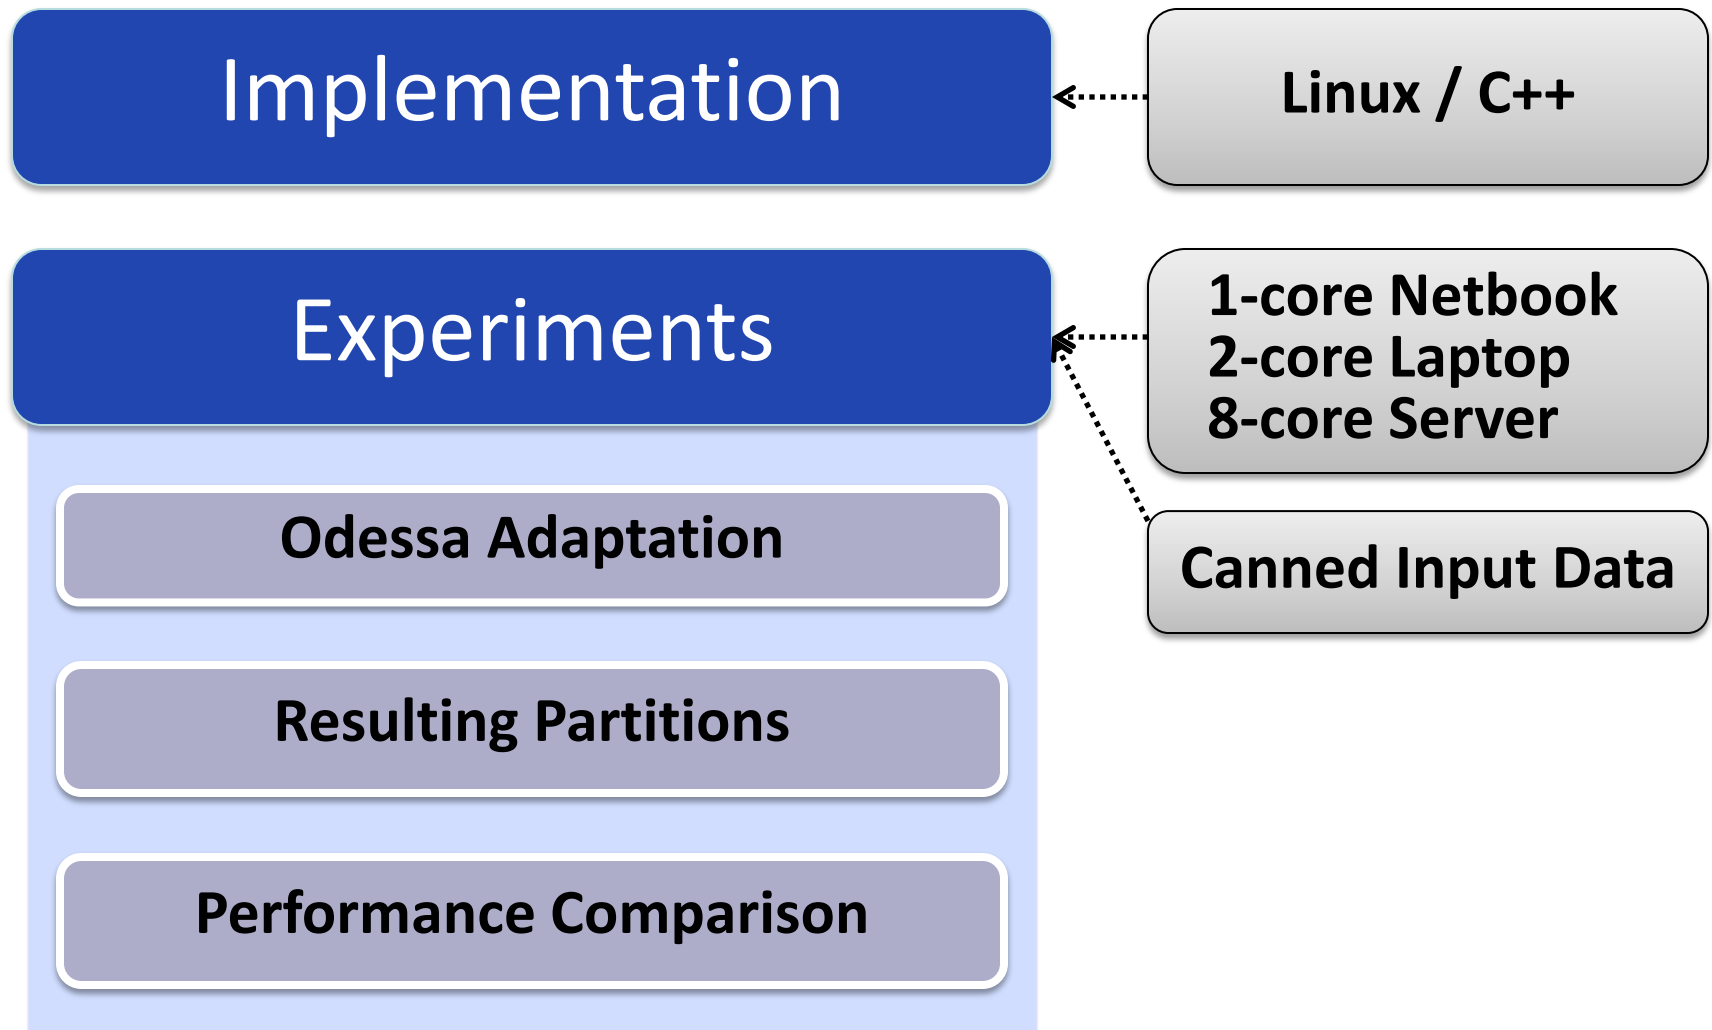
Cloud Infrastructure




Incremental decisions adapt quickly to input and platform variability.

Smartphone

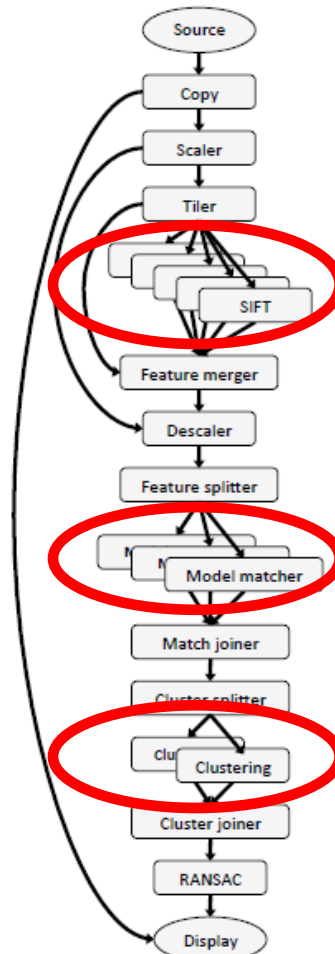
Evaluation Methodology



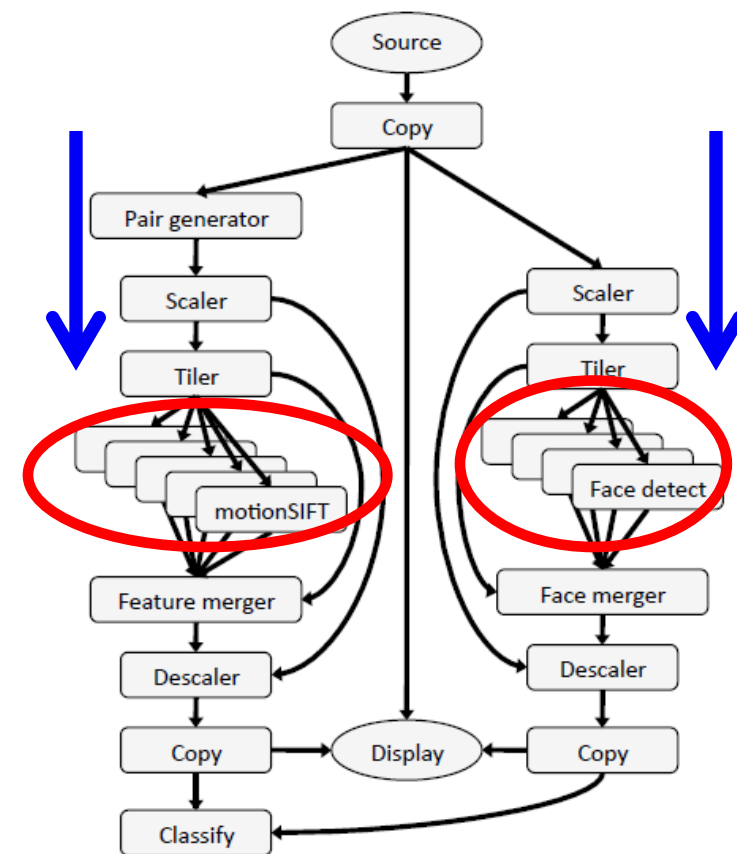
Data-Flow Graph



Face Recognition



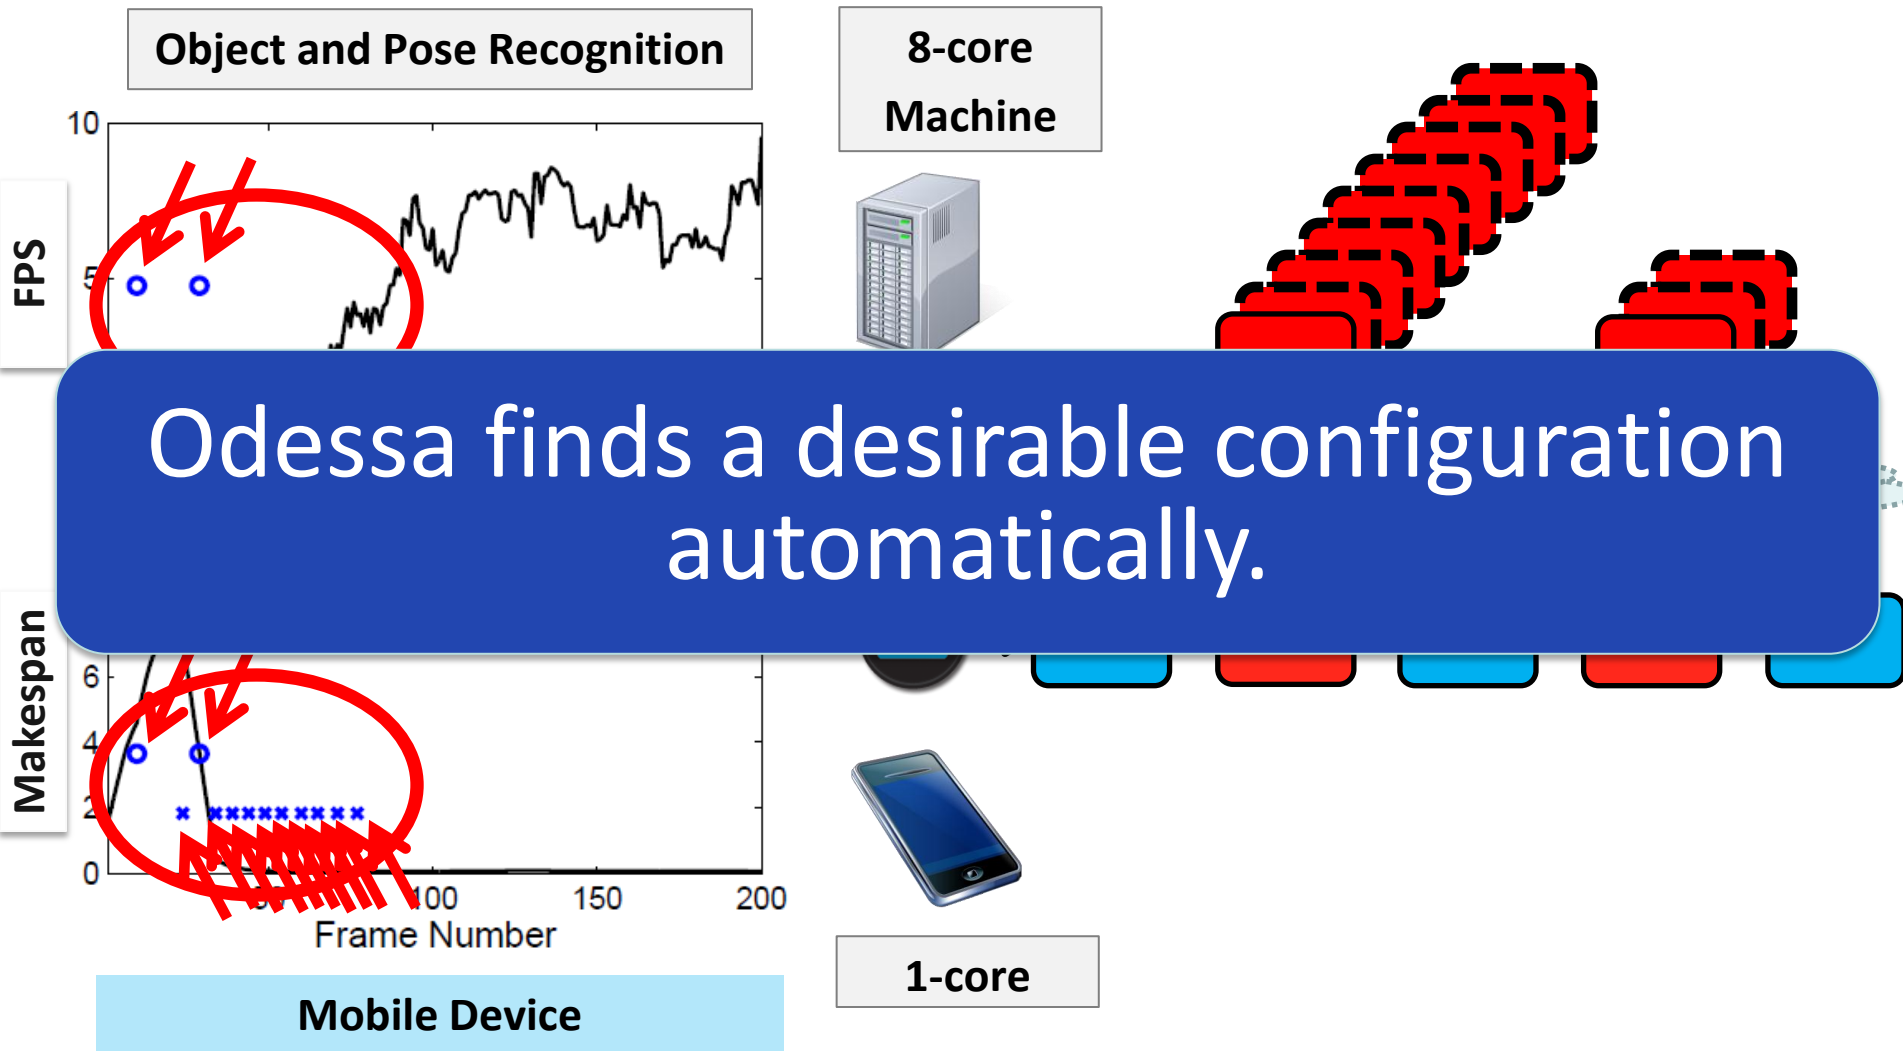
Object Pose Estimation



Gesture Recognition



Odessa Adaptation



Resulting Partitions in Different Devices

Face Recognition

Client Device	Stage Offloaded and Instances	Degree of Pipeline Parallelism
Mobile Device	Face detection (2)	3.39

Resulting partitions are often very different for different client devices.

Client Device	Stage Offloaded and Instances	Degree of Pipeline Parallelism
Mobile Device	Face Detection (1) Motion-SIFT Feature (4)	3.06
Dual Core Notebook	Face Detection (1) Motion-SIFT Feature (9)	5.14



Performance Comparison with Other Strategy

Object and Pose Recognition Application

Strategy	Throughput (FPS)	Makespan (Latency)
Odessa performs 4x better than the partition suggested by domain expert, close to the offline optimal strategy.		
Offline-Optimal	6.49	430 ms
Odessa	6.27	807 ms

Mobile Device



Related Work

- *ILP solver* for saving energy: [MAUI] [CloneCloud]
- *Graph-based* partitioning: [Gu'04] [Li'02] [Pillai'09] [Coign]
- *Static Partitioning*: [Wishbone] [Coign]
- A set of *pre-specified* partitions: [CloneCloud] [Chroma] [Spectra]

Objectives

Variability

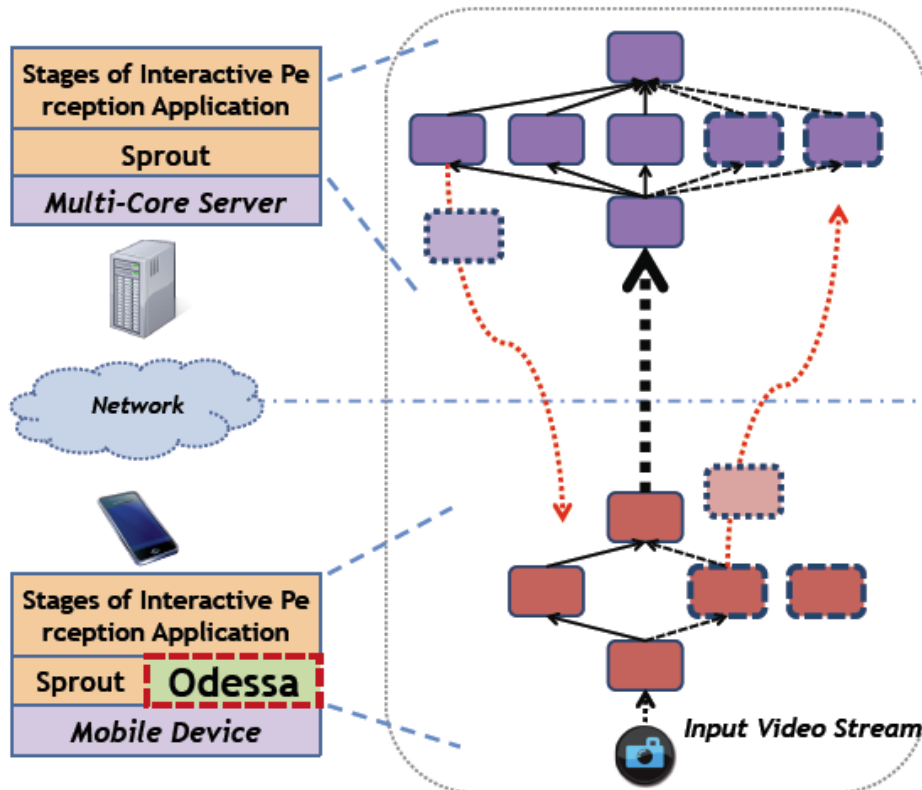
Migration,
Contention

Parallelization

Odessa



Summary of Odessa



Adaptive & Incremental runtime for mobile perception applications

- Odessa system design using novel workloads.
- Understanding of the factors which contribute to the offloading and parallelism decisions.
- Extensive evaluation on prototype implementation.

Thank you

“Any questions?”